



DEPARTMENT OF
NETWORKED SYSTEMS
AND SERVICES

COMPUTER ARCHITECTURES

Instruction set architectures

Gábor Horváth, ghorvath@hit.bme.hu


Budapest,
04/27/2023



- Every processor has
 - Instruction Set Architecture (ISA)
 - A kind of programming interface
- Parts of ISA:
 - List of supported instructions
 - List of supported data types
 - Registers
 - Addressing modes
 - Flags
 - How to communicate with I/O devices
 - Interrupt and exception handling
 - Etc.

- „Parts” of an instruction:

- Code of the operation
- Values/addresses of operands
- Address where the result is stored
- Pointer to the next instruction



- For simplicity:

- Pointer to the next instruction is unnecessary
 - The next instruction is always the next one in the memory
- Number of operands
 - Support for 3 operands: **R1 ← R2 + R3**
 - Support for 2 operands: **R1 ← R1 + R2**
 - Support for 1 operands: **ADD R1**

- Instruction types:

- Data movement

$R1 \leftarrow R2$, $R1 \leftarrow \text{MEM}[100]$, $R1 \leftarrow 42$, $\text{MEM}[100] \leftarrow R1$, $\text{MEM}[100] \leftarrow 42$

- Arithmetic-logic operations

$R1 \leftarrow R2 + R3$, $R1 \leftarrow \text{MEM}[100] * 42$, $\text{MEM}[100] \leftarrow R1 \ \& \ R2$

- Control flow operations:

$\text{JUMP } -42$, $\text{JUMP } +28 \text{ IF } R1 == R2$, CALL proc , RETURN

- Stack operations

$\text{PUSH } R1$, $\text{PUSH } 42$, $R2 \leftarrow \text{POP}$

- I/O operations

$\text{IO}[42] \leftarrow R1$, $R1 \leftarrow \text{IO}[42]$

- Transcendent operations

$R2 \leftarrow \text{SIN } R1$, $R2 \leftarrow \text{SQRT } 42$

- Etc.

- Determines, where the operand is located
- Possible locations:
 - embedded in the instruction (immediate)
 - in a register
 - in the memory

Addressing mode	Example
Register	$R1 \leftarrow R2 + R3$
Immediate	$R1 \leftarrow R2 + 42$
Direct	$R1 \leftarrow R2 + \text{MEM}[42]$
Register indirect	$R1 \leftarrow R2 + \text{MEM}[R3]$
Indirect with offset	$R1 \leftarrow R2 + \text{MEM}[R3+42]$
Memory indirect	$R1 \leftarrow R2 + \text{MEM}[\text{MEM}[R3]]$
Indexed	$R1 \leftarrow R2 + \text{MEM}[R3+R4]$

- Typically relative addressing, e.g. JUMP -28
- 3 possible implementations of conditional branches:
 - With condition codes:
COMPARE R1, R2
JUMP label IF GREATER
 - With condition registers:
R0 ← R1 > R2
JUMP label IF R0
 - „Compare and jump”:
JUMP label IF R1 > R2



- Predicates: instructions with conditions

R1 ← R2 + 32 IF P2

- Predicate registers: 1-bit registers (**P2** in the example)
- Setting the predicate registers: by comparison instructions

P2 ← R3 ≤ R5

- The instructions are stored with a binary encoding
- Based on the length of the encoded instructions we have:
 - Fixed length encoding
 - Variable length encoding



- Low-level programming:
 - Manual encoding of the instructions is inconvenient
 - Binary coded instructions are not for human use
 - Tool: assembly programming
- Assembly
 - The lowest level programming language
 - The text representation of the machine level instructions
 - 1 assembly „instruction” → 1 machine level instruction
- Assembler: creates machine code from assembly
- It is different for every instruction set architecture

ENCODING THE INSTRUCTIONS



ENCODING THE INSTRUCTIONS



- The CPU of the Terminator: MOS-6502 (like the CPU of Apple II) (the cheapest CPU between 1975 and 1980: it costed sixth the price of the Intel and Motorola CPU-s at the same level of performance)
- The Terminator runs an example code of the Nibble magazine

ENCODING THE INSTRUCTIONS

Assembly code

Binary encoded
(machine) instructions

```

136 6085 A0 0A
137 6087 B9 E1 60
138 608A 99 F0 03
139 608D 8C
140 608E 10 F7
141 6090 AD F3 03
142 6093 49 A5
143 6095 8D F4 03
144
145
146
147 6098 A9 03
148 609A A0 F0
149 609C 85 3D
150 609E 85 43
151 60A0 84 3C
152 60A2 84 42
153 60A4 A9 03
154 60A6 A0 FA
155 60A8 85 3F
156 60AA 84 3E
157 60AC 38
158 60AD 20 11 C3
159
160
161
162 60B0 A9 60
163 60B2 A0 EF
164 60B4 85 3D
165 60B6 85 43
166 60B8 84 3C
167 60BA 84 42
168 60BC A9 64
169 60BE A0 5D
170 60C0 85 3F
171 60C2 84 3E
172 60C4 38
    
```

```

2833
2848
2A5C
295C
2A09
21C0
243D
28FD
279F
BB
PROGRAM CHECK 1S : 0502
9310 - 935F
9360 - 93AF
93B0 - 93FF
9400 - 944F
9450 - 949F
94A0 - 94EF
94F0 - 953F
9540 - 958F
9590 - 95DF
95E0 - 95E1
    
```

```

INSTALL
LOOP
LDY #10
LDA VCT, Y
STA PG3VEC, Y
DEY
BPL LOOP
LDA $3F3
EOR $A5
STA $3F4

and aux mem

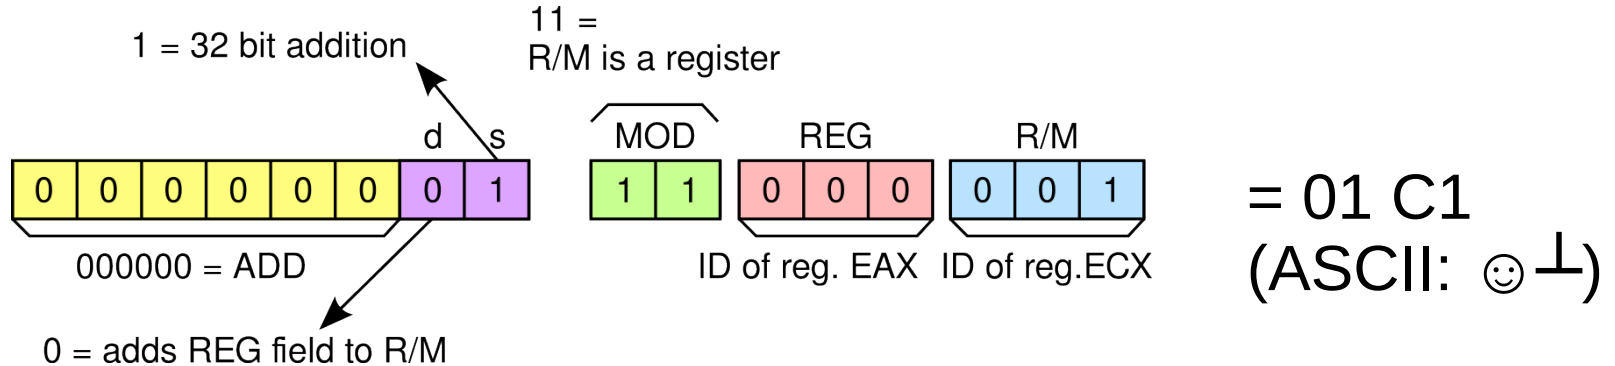
LDA #PG3VEC/
LDY #PG3VEC
STA A1H
STA A4H
STY A1L
STY A4L
LDA #PG3VEC+10/
LDY #PG3VEC+10
STA A2H
STY A2L
SEC
JSR AUXMOVE

Copy program to aux memory

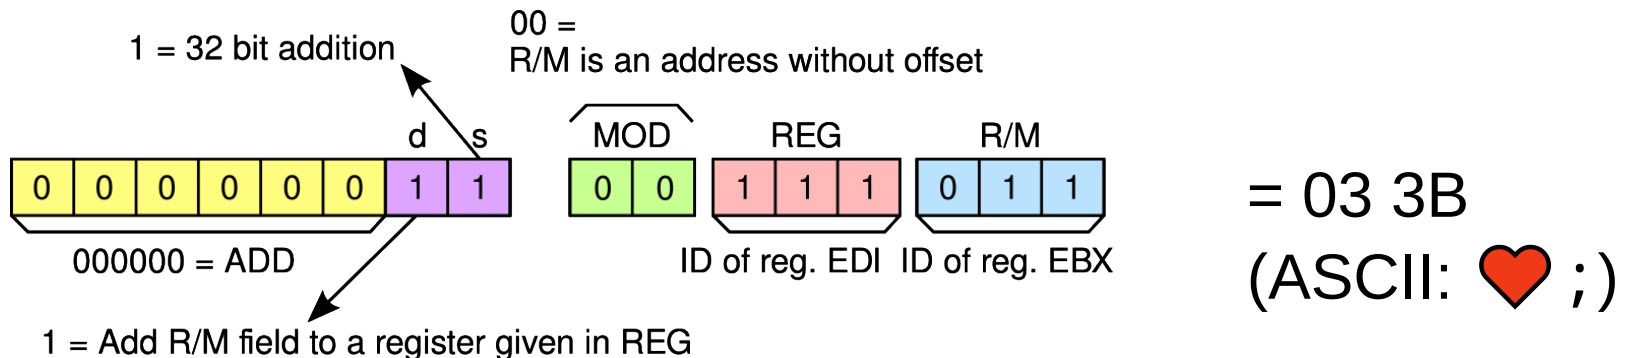
LDA #BEGIN/
LDY #BEGIN
STA A1H
STA A4H
STY A1L
STY A4L
LDA #END/
LDY #END
    
```

EXAMPLE FOR INSTRUCTION CODING (X86)

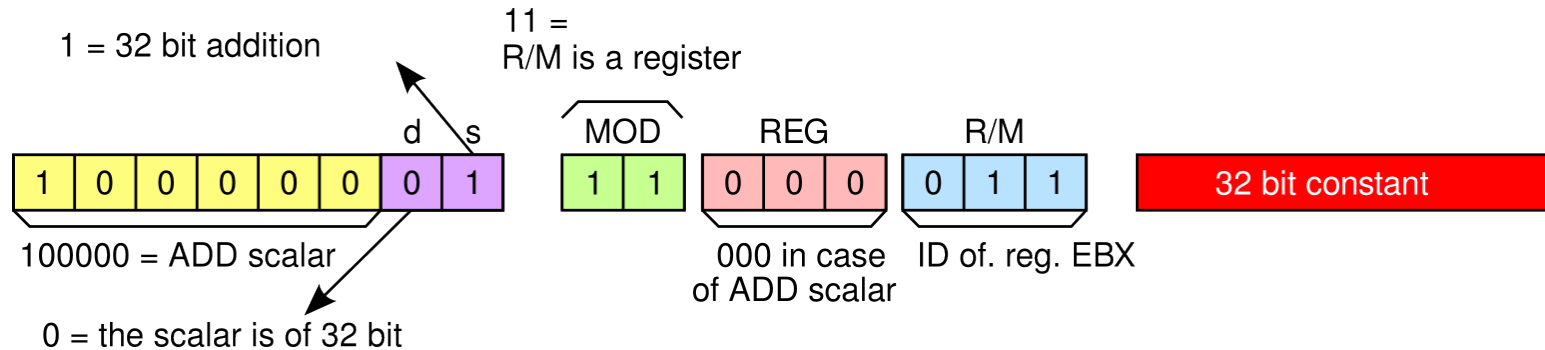
- **ADD ECX, EAX** (we write it: $ECX \leftarrow ECX + EAX$)



- **ADD EDI, [EBX]** (we write it: $EDI \leftarrow EDI + MEM[EBX]$)



- **ADD EBX, 23423765** (we write it: $EBX \leftarrow EBX + 23423765$)



- **= 81 C3 15 6B 65 01** (ASCII: Qüşke ☺)

- Byte order:
 - Little endian: numbers start with the **least** significant byte
 - Big endian: numbers start with the **most** significant byte
 - Bi endian: can be selected (HW or SW)
 - Example: 23423765 (=1656B15)
 - Little endian: 15 6B 65 01
 - Big endian: 01 65 6B 15
- Ways of communicating with I/O devices:
 - Separate instructions for I/O
 - Memory mapped

- Typical in the 70's:
 - a large number of instructions
 - complex instructions
- Motivations:
 - Memory was slow – with complex instructions, the processor got more work with a single memory operation
 - Memory was expensive – a single instruction describes more work for the processor
 - Compilers were very basic that time. The processor had a „high level” like instruction set to allow easy assembly programming.
- This is the **CISC** (Complex Instruction Set Computer) philosophy
- Features:
 - Easy to use instructions
 - Register-memory instructions (pl. $R1 \leftarrow R2 + \text{MEM}[42]$)
 - Redundancy
 - Several addressing modes
 - Variable length instruction encoding
 - The execution time of the various instructions is different

- Typical CPU design trends in the 80's and 90's:
 - The number of instructions is as small as possible
 - The instructions are very basic
- Motivations:
 - To simplify the design of the CPU-s
 - The simpler CPU design allows more efficient implementation
- This is the **RISC** (Reduced Instruction Set Computing) philosophy
- Features:
 - Simple, elementary instructions, avoiding redundancy
 - Load-Store and register-register operations
 - instead of **$R1 \leftarrow R2 + \text{MEM}[42]$** we have **$R3 \leftarrow \text{MEM}[42] ; R1 \leftarrow R2 + R3$** .
 - Only a few addressing modes are available
 - Fixed instruction encoding
 - Execution time of the instructions usually takes only 1 cycle

- Comparison:
 - CISC: dense
 - The program is smaller
 - RISC: simple
 - Less design bugs
 - The IC is smaller
 - It consumes less energy
 - Better yield when manufacturing it
 - There is a lot of space left on the IC allowing the integration of further devices onto the same silicon
 - CISC: a small number of registers vs. RISC: much more registers



Some important instruction set architectures

x86

- First appearance: 1971, Intel 8086
- 1981: The Intel 8088 is selected as the CPU of the IBM PC
- Originally it was a 16 bit ISA, but has been extended to 32 and 64 bit later
- Nowadays it is used both in high-performance servers and low-power mobile devices
- A very obsolete ISA, but the demand for software compatibility keeps it alive for >50 years
- Intel has spent most of its profit to develop more efficient semiconductor production technology

ARM

- First implementation: 1987
- The most wide-spread ISA all over the world
- It is 32 bit right from the beginning (extension to 64 bit is in the works)
- This ISA is very carefully designed, easy to implement
 - Can be implemented with only 30.000 transistors!
- ARM does not manufacture CPUs
 - The ISA can be licensed
 - ARM designs CPUs, that can be licensed as well (ARM Cortex family)
- Primary goals: simplicity, energy efficiency
(**not** the raw computational power)

POWER

- Defined in 1991 by IBM, Apple and Motorola
- Goal: to surpass the computational performance of x86
- They succeeded:
 - Huge memory and I/O bandwidth
 - 2014: 5 GHz, 12 cores, 8 threads/core (POWER8)
- Did not get popular in PCs
- But it got popular in workstations and servers
- ... and all prev. gen. game consols used POWER processors!
(Microsoft XBox 360, Sony PlayStation 3, Nintendo Wii)

SPARC

- 1987, SUN
- 64 bit from the beginning
- Open platform!
- The design of UltraSPARC T1 and T2 can be accessed by anybody (at VeriLog level)
- Still in production (now by Oracle)
 - 2013: SPARC T5: 16 cores, 8 threads/core, 3.6 GHz, etc.
 - 2016: SPARC M7: 4.13 GHz, 32 cores, 8 threads/core (256 threads!)
- In 2017, the 7th most powerful computer is SPARC based (and it is the first without GPU)

Alpha (DEC, 1992)

- 64 bit from the beginning
- Extremely innovative:
 - 21164: the first CPU with a big cache integrated with the CPU
 - 21264: the first CPU with both high frequency and out-of-order execution
 - 21364: the first CPU with integrated memory controller
 - 21464: supposed to be the first multi-thread CPU (but the project was stopped meanwhile)
- Extremely strong floating point unit
- 21264 @ 833 MHz > 3x Pentium III @ 1 GHz!
- Hand-made design
- Canceled when Compaq acquired DEC

PA-RISC (1986, HP)

- First 32, later 64 bit CPUs
- Extremely strong floating point unit
- PA-8600 @ 552 MHz > 2x Pentium III @ 1 GHz!
- Canceled when HP started to develop the Itanium processors with Intel

IA-64 (Itanium)

- 1994, joint development of HP & Intel
- Huge interest from the press, very costly development
- First implementation: 2001, disappointing performance, sold only few thousand
- Supposed to be compatible with x86: succeeded, but it can reach only the level of a Pentium clocked at 100MHz...
- Problem: it needs a special compiler to utilize its abilities, they did not count with the difficulties of developing such compiler
- Still developed and manufactured, sold only 55.000 between 2001-2007
- Most big companies stopped supporting it
 - 2008: Microsoft
 - 2011: Oracle
- 2018: Intel announced to finish the production (till 2021)

	x86	ARM	PowerPC	SPARC
Number of bits	64	32	64	64
Year introduced	1978	1983	1991	1985
Num of operands	2	3	3	3
Instruction style	Reg-mem	Reg-reg	Reg-reg	Reg-reg
CISC vs. RISC	CISC	RISC	RISC	RISC
Num of registers	8/16	16	32	32
Instruction coding	Variable (1-17)	Fixed (4)	Fixed (4 – com.)	Fixed (4)
Conditional instr.	Condition code	Condition code	Condition code	Condition code
Byte order	Little	Big	Big/Bi	Bi
Addressing modes	5	6	4	2
Branch predication	No	Yes	No	No

	m68k	Alpha	PA-RISC	Itanium
Number of bits	32	64	64	64
Year introduced	1979	1992	1986	2001
Num of operands	2	3	3	3
Instruction style	Reg-mem	Reg-reg	Reg-reg	Reg-reg
CISC vs. RISC	CISC	RISC	RISC	EPIC
Num of registers	16	32	32	128
Instruction coding	Variable (2-22)	Fixed (4)	Fixed (4)	Fixed (16)
Conditional instr.	Condition code	Condition reg.	Compare & jump	?
Byte order	Big	Bi	Big	Bi
Addressing modes	9	1	5	?
Branch predication	No	No	No	Yes



DEPARTMENT OF
NETWORKED SYSTEMS
AND SERVICES

