



DEPARTMENT OF
NETWORKED SYSTEMS
AND SERVICES


COMPUTER ARCHITECTURES

Instruction Pipeline

Gábor Horváth, ghorvath@hit.bme.hu

Budapest,
2023. 05. 02.





- Our CPU in the examples: RISC
- Instructions:
 - Load / Store
 - $R1 \leftarrow \text{MEM}[R0+42]$ or $\text{MEM}[R0+42] \leftarrow R1$
 - Addressing mode: indirect with offset
 - Arithmetic / Logic
 - $R1 \leftarrow R2 + R3$ or $R1 \leftarrow 42 * R3$ or $R4 \leftarrow R1 \& R5$
 - Control
 - Branch, jump
 - Conditional jump
 - $\text{JUMP } -24 \text{ IF } R2 == 0$
- Only Load and Store are allowed to access the memory

- Stages of the execution of instructions:

- 1) Instruction Fetch, **IF**

- + increase program counter (if it is not a jump)

- 2) Instruction decode/register fetch, **ID**

- Decoding the binary encoded instruction
 - Determine what kind of instruction it is
 - also provides control signals for ALU
 - Operands located in registers are retrieved from the register file

- 3) Execution, **EX** – ALU executes an operation

In case of

- Load/Store: it calculates the memory address
(**R1** ← **MEM**[**R0** + **42**])
- Arithmetic/Logic: it calculates the result (**R1** ← **R2** * **R3**)
- Conditional jump: it evaluates the condition and also calculates branch target address (**JUMP** [**PC**] - **24** **IF** **R2** == **0**)

... cont'ed:

4) Memory access, **MEM**

In case of

- Load/Store: the memory operation is accomplished
- Arithmetic/logic, jumps: this phase does nothing


5) Write-back, **WB**

In case of


- Arithmetic/logic: stores the result to the register file
(**R1** \leftarrow **R2*****R3**)
- Load: stores the memory content retrieved to register file
(**R1** \leftarrow **MEM**[**R0**+**42**])
- Store, jumps: this phase does nothing

- Stages to execute:
 - Arithmetic instructions: IF, ID, EX, WB
 - Store instructions: IF, ID, EX, MEM
 - Load instructions: IF, ID, EX, MEM, WB
 - Jumps: IF, ID, EX
 - Etc.


EXECUTING INSTRUCTIONS



EXECUTING INSTRUCTIONS




- This was the concept of instruction pipeline
- Stages of instruction executions are overlapped:
 - Every stage needs to be given the same amount of time:
cycle time

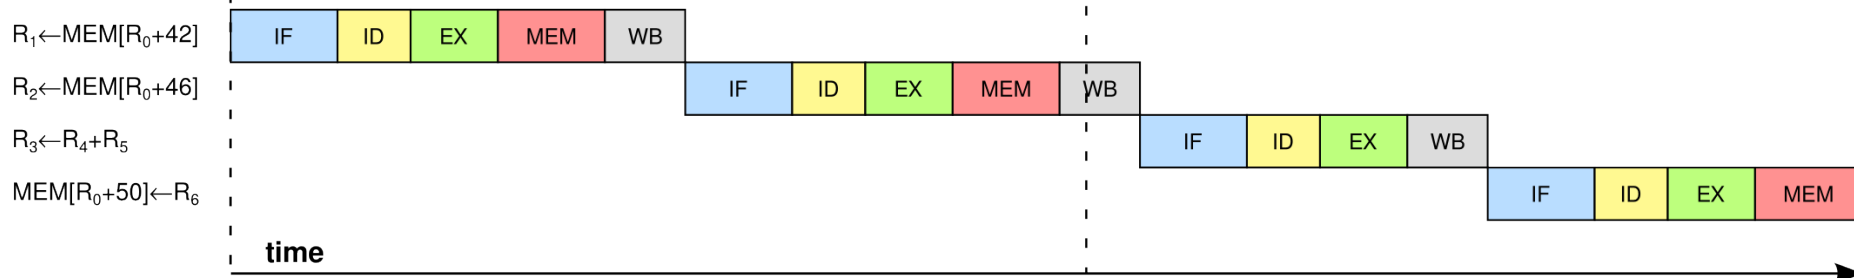


- How much do we gain?
 - On the one hand: we overlapped the execution of several instructions → gain
 - On the other hand: there are idle time intervals (some stages are not needed by all types of instructions, some stages are faster than the cycle time, etc.)
- As long as the instructions are independent, the gain is significant

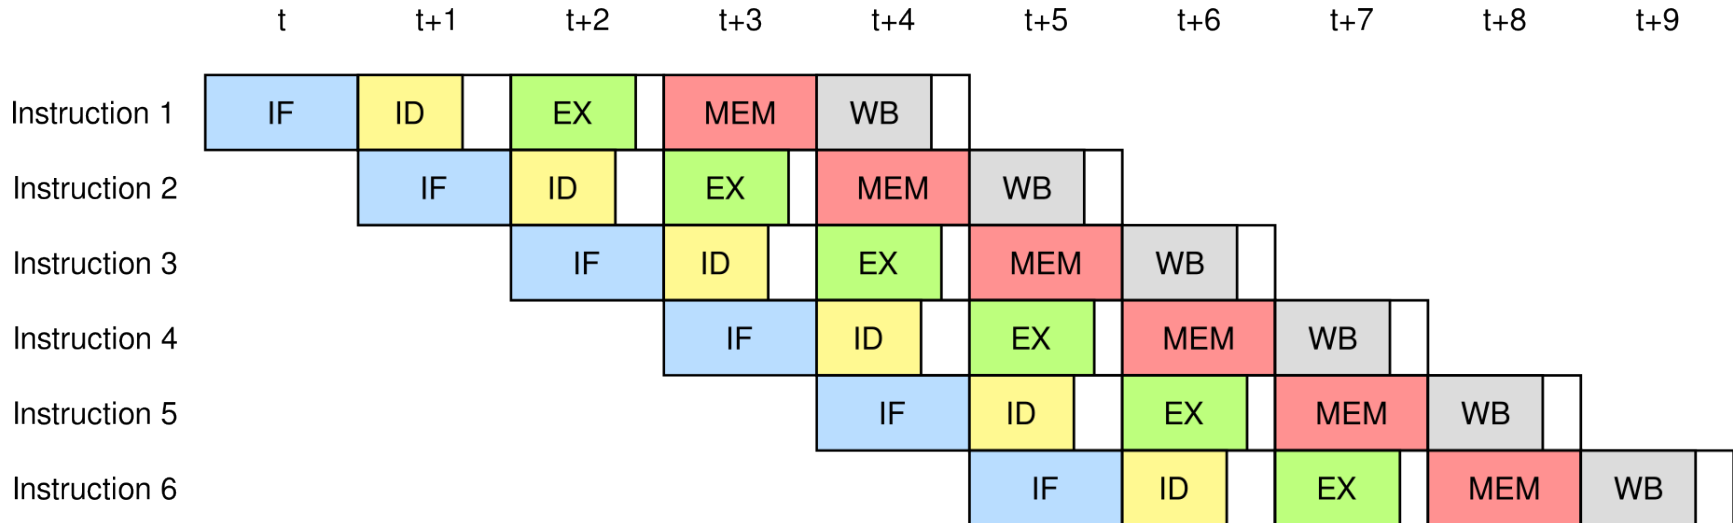
Using pipeline:




Without pipeline:



- Quantities:
 - Depth: number of stages
 - Latency: the execution time of an instruction
 - Throughput: executed instructions / sec



- Each stage needs to exchange some information with the neighboring stages about the current instruction
→ these are stored in **pipeline registers**



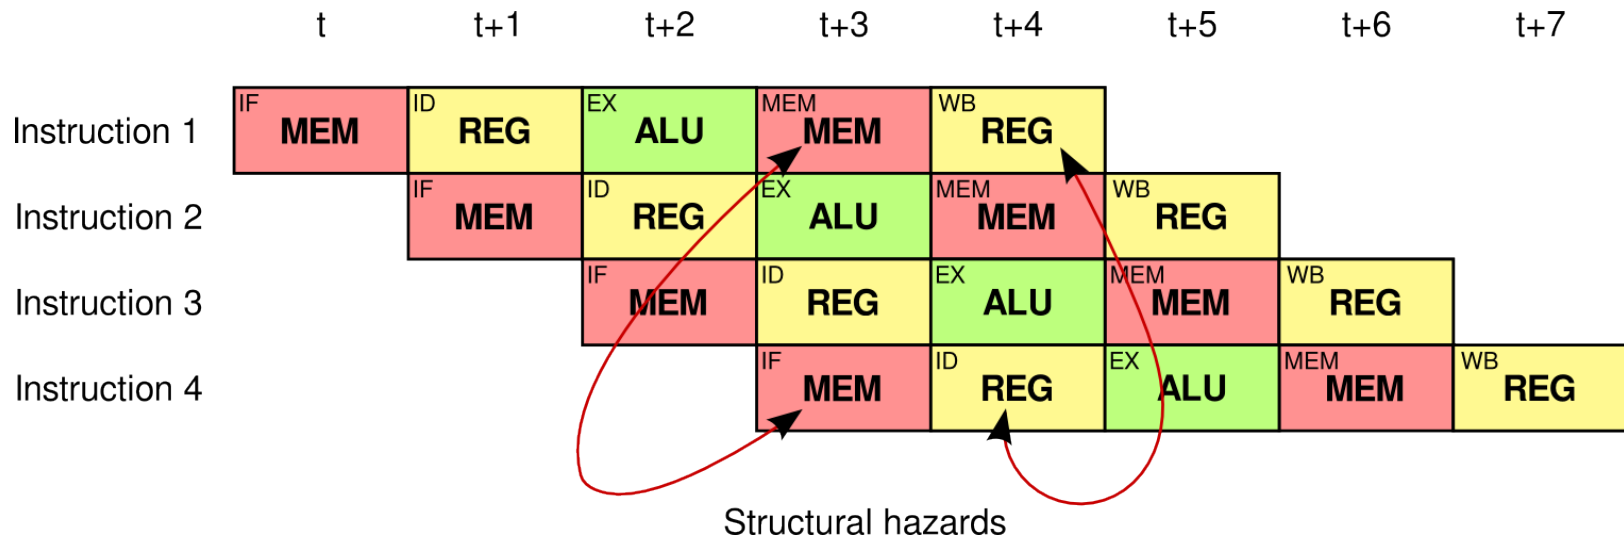
- Example:
 - IF puts the fetched instruction to register IF/ID
 - ID takes it out from IF/ID, and puts the ALU control signals and operands to register ID/EX
 - EX takes it out from ID/EX, executes the arithmetic operation, puts the result to register EX/MEM
 - Etc... (later)

- There is a serious problem
- The instructions are not independent!
 - Sometimes they need the same system resources at the same time
→ **Structural hazard**
 - An operand of an instruction is the result of the previous instruction
→ **Data hazard**
 - Conditional jumps: we do not know where the execution of the program continues until the branch condition is evaluated
→ **Control hazard**

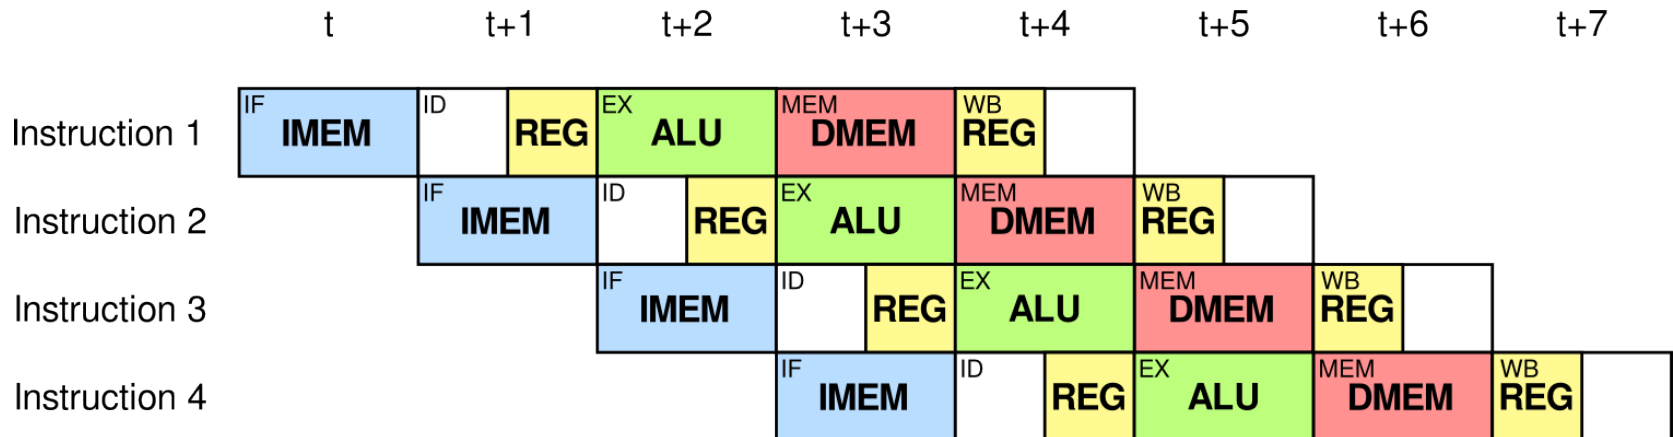
How to treat them:

- We either resolve them by using some smart trick
- ...or stop the pipeline till the hazard persists
→ degrades efficiency

- Some system resources are used in several stages as well
 - IF: **Memory**
 - ID: **Register file**
 - EX: **ALU**
 - MEM: **Memory**
 - WB: **Register file**



- Resolving it:
 - In case of the „memory” resource:
 - By using separate instruction and data cache
 - In case of the „register file” resource:
 - One of the stages uses the register file in the first half of the cycle time, the other stage uses it in the second half



- Caused by data dependencies
- Data hazard:

Multiple instructions operate on the same data address
(on the same register, or on the same memory address)

- E.g. an operand of an instruction is the result of a previous instruction


→ **RAW dependency** (Read After Write)

```
i1:  R3 ← MEM[R2]
i2:  R1 ← R2 * R3
i3:  R4 ← R1 + R5
i4:  R5 ← R6 + R7
i5:  R1 ← R8 + R9
```

- RAW dependencies:

i1 ↔ i2, i2 ↔ i3

RAW DEPENDENCIES

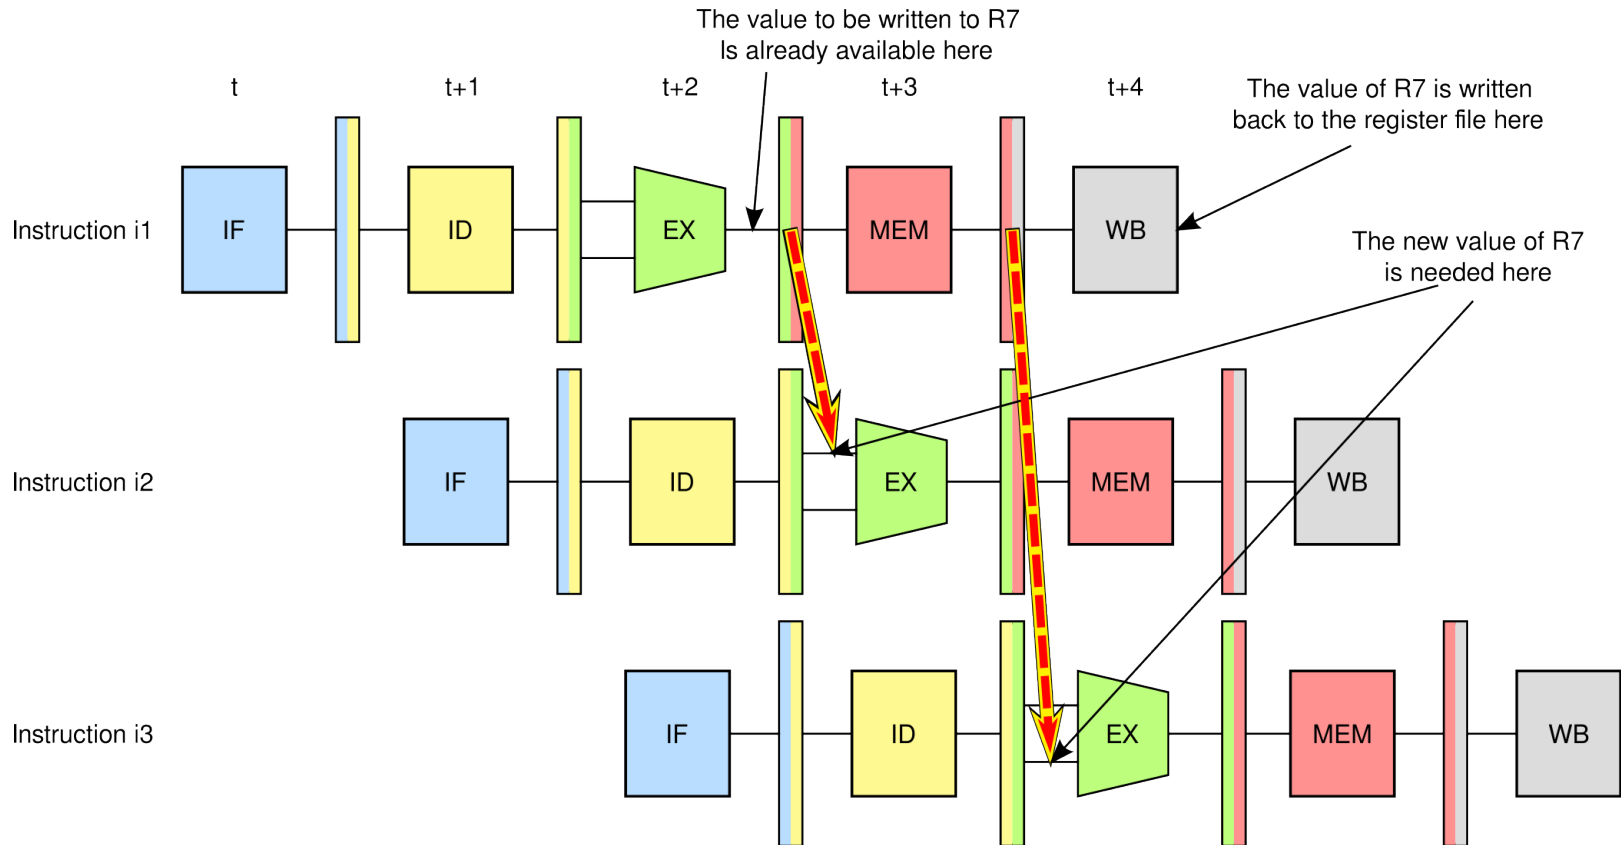


i1: R7 ← R1 + R5

i2: R8 ← R7 + R2

i3: R5 ← R8 + R7

- The result is available, but it is not yet stored to the register file
- Let us read it out from the appropriate pipeline register → **forwarding**



i1: $R7 \leftarrow R1 + R5$

i2: $R8 \leftarrow R7 + R2$

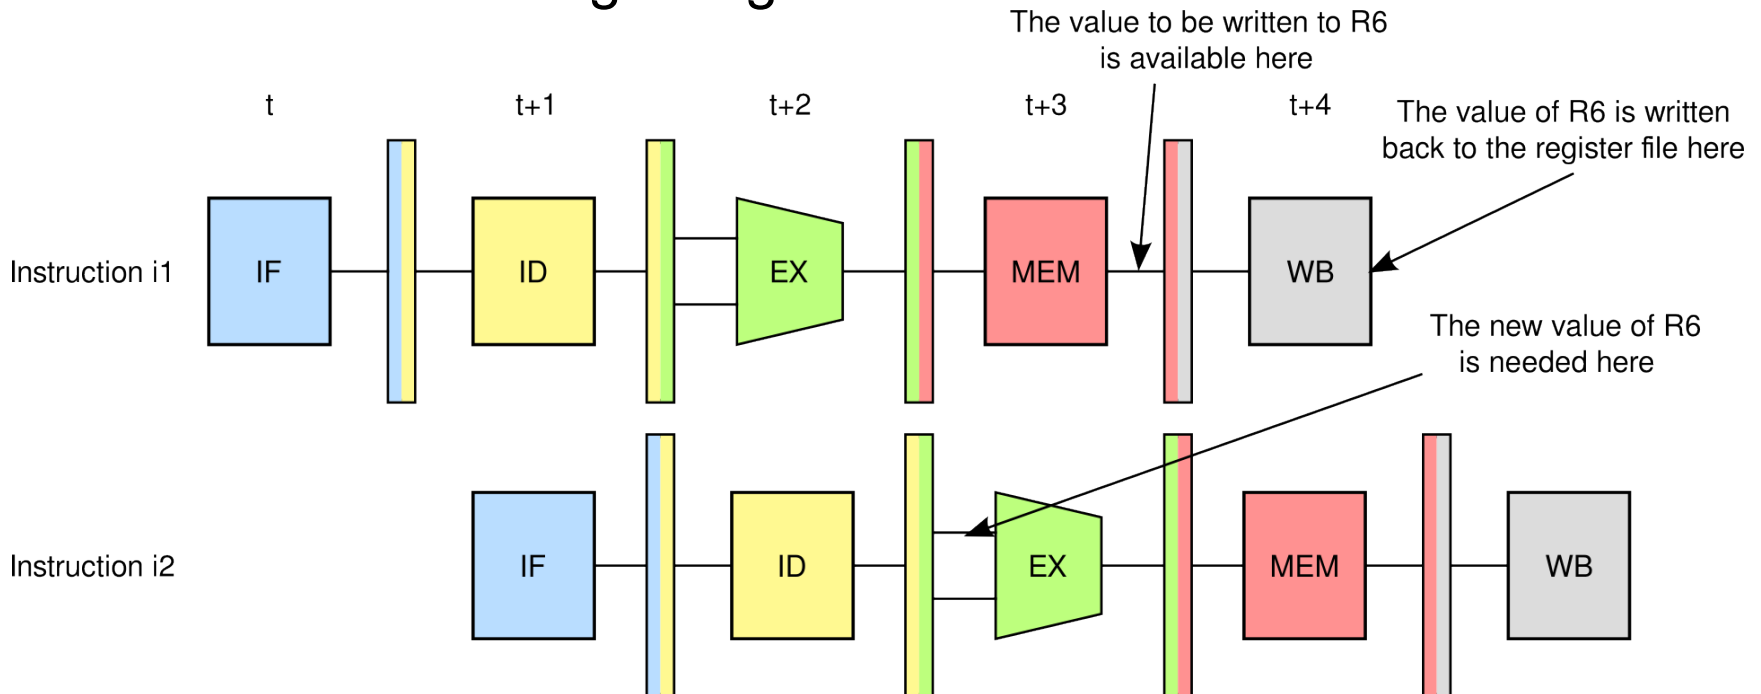
i3: $R5 \leftarrow R8 + R7$

- Forwarding: there are situations when it does not help

i1: R6 ← MEM[R2]

i2: R7 ← R6 + R4

- The value of R6 is available at the end of MEM stage
- i2 needs it at the beginning of EX

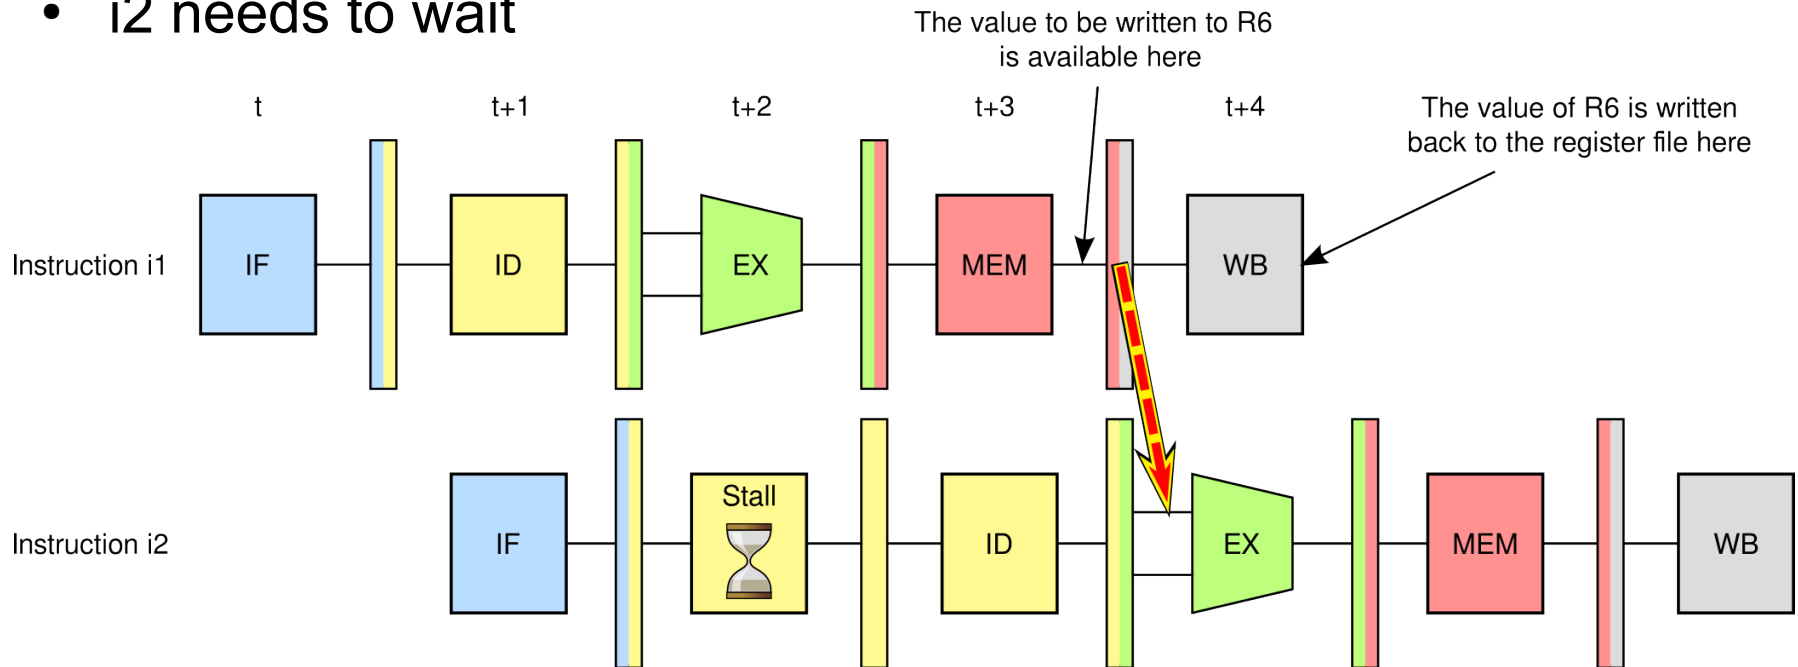


- Forwarding: there are situations when it does not help

i1: R6 ← MEM[R2]

i2: R7 ← R6 + R4

- The value of R6 is available at the end of MEM stage
- i2 needs it at the beginning of EX
- i2 needs to wait



- Data dependencies:
 - **RAW**: an instruction reads the register/address written by an earlier instruction
 - Solution: forwarding/waiting
 - **WAR**: an instruction writes the register/address read by an earlier instruction
 - Not a problem in this simple pipeline
 - **WAW**: multiple instructions write the same register/address
 - Not a problem in this simple pipeline
 - **RAR**: multiple instructions read the same register/address
 - Never a problem

- Example:

```
i1:  R3 ← MEM[R2]
i2:  R1 ← R2 * R3
i3:  R4 ← R1 + R5
i4:  R5 ← R6 + R7
i5:  R1 ← R8 + R9
```

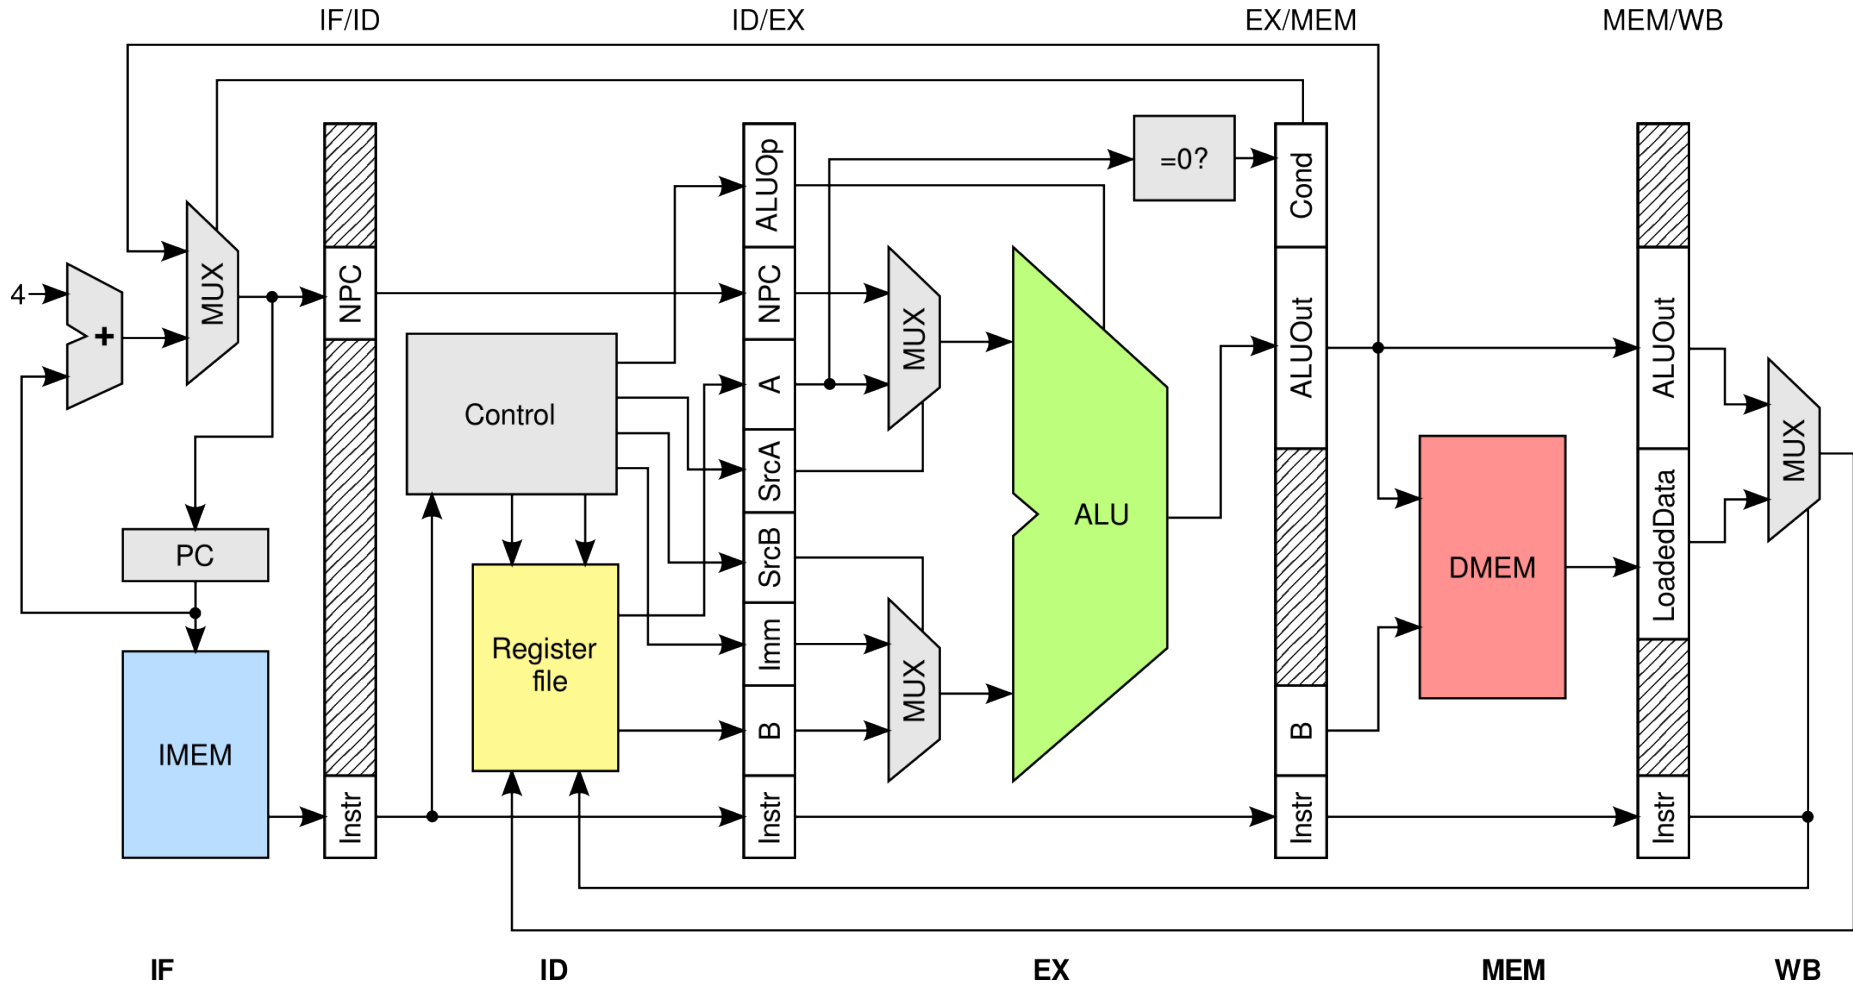
- **WAW**: i2 ↔ i5 (write after write)
- **WAR**: i3 ↔ i4 (write after read)

- There is a problem with the conditional jump instructions:
 - IF fetches them
 - The jump condition and the jump target address is calculated only by stage EX!
 - Where to fetch the subsequent instruction from till it completes the EX stage?
- Solutions:
 - Stop and wait till the jump instruction completes the EX stage
→ used by Intel 80386
 - Predict the jump condition and the address immediately
→ ***branch prediction***

- **Static branch prediction** → not adaptive
There are methods which
 - Always predicts that the jump won't be taken
 - It costs nothing
 - In case of misprediction, the instructions fetched by mistake are invalidated (by Intel 80486)
 - However, conditional jumps are more often taken than not taken (loops)
 - Always predicts that the jump will be taken
 - OK, but how do we know the jump target address then?
 - It works well if the target address is calculated well before the condition
 - Jump forward: not taken, jump backward: taken
 - Motivation: loops usually do jump back, except, when they are finished
- **Dynamic branch prediction**
 - Adapts the behavior of the program (it learns which jumps are taken and which are not)
 - It is much better than the static one
 - All modern CPUs use dynamic branch prediction

- What did we learn up to this point?
 - The concept of instruction pipeline
 - That the pipeline stages communicate by using the pipeline registers
 - How to cope with hazards
- Let us design our own instruction pipeline!

IMPLEMENTATION OF THE PIPELINE



- ALU: has a central role
- It does computations:

R1 \leftarrow **R1** * **42**

R1 \leftarrow **R2** & **R3**

PC \leftarrow **PC** + **42** (JUMP +42)


R1 \leftarrow **MEM** [**R2** + **42**]

- First operand: **register** or **PC**
- Second operand: **register** or **constant** (immediate)
- Evaluates jump conditions:


JUMP -28 IF **R1==0**

- The task of the ID phase:
 - Preparing operands for the ALU
 - Preparing operand selection signals for the ALU
 - Preparing the operation code (+, -, *, /, &, etc.) for the ALU


- Updates the program counter:
 - If jump (**EX/MEM.Instr Opcode==branch**),
and if it will be taken (**EX/MEM.Cond==TRUE**):
 - **PC \leftarrow EX.MEM.ALUOut**
 - Otherwise:
 - **PC \leftarrow PC+4**
- Passing the instruction word and the new program counter value forward:
 - **IF/ID.NPC \leftarrow PC**
 - **IF/ID.Instr \leftarrow IMEM[PC]**



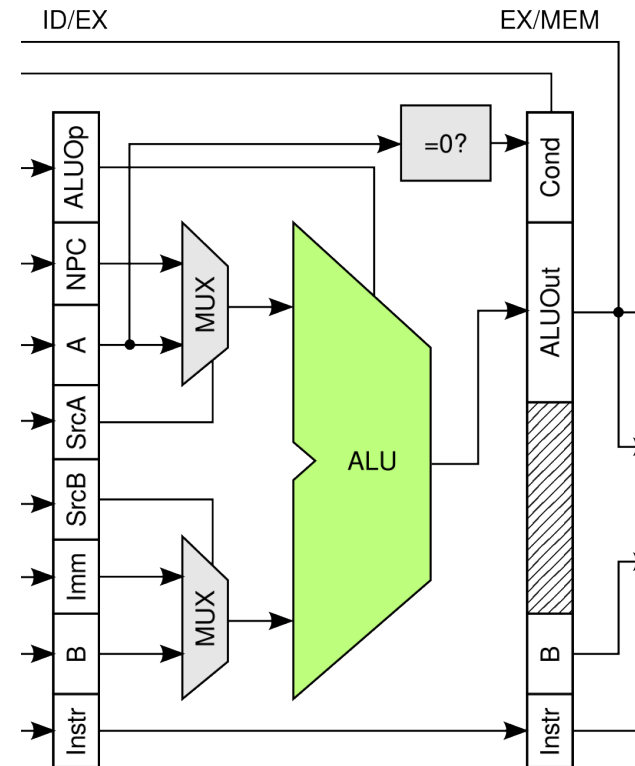
- Prepares the operands:
 - First operand:
 - $ID/EX.NPC \leftarrow IF/ID.NPC$
 - $ID/EX.A \leftarrow Reg [IF/ID.Instr.ra]$
 - Selection signal for the first operand:
 - In case of jump instruction ($IF/ID.Instr.Opcode == branch$):
 - $ID/EX.SrcA \leftarrow npc$
 - Otherwise:
 - $ID/EX.SrcA \leftarrow regA$
 - Second operand:
 - $ID/EX.Imm \leftarrow IF/ID.Instr.imm$
 - $ID/EX.B \leftarrow Reg [IF/ID.Instr.rb]$
 - Selection signal for the second operand:
 - If constant ($IF/ID.Instr.HasImm$):
 - $ID/EX.SrcB \leftarrow imm$
 - Otherwise:
 - $ID/EX.SrcB \leftarrow regB$
 - ...




- ...
- Prepares operation code:
 - If we have an arithmetic instruction
(**IF/ID.Instr.Opcode == arithm**):
 - **ID/EX.ALUOp ← IF/ID.Instr.Func**
 - Otherwise (PC update or address calculation):
 - **ID/EX.ALUOp ← „+”**
- Passing the instruction word forward:
 - **ID/EX.Instr ← IF/ID.Instr**




- ALU inputs:
 - If **ID/EX.SrcA == npc**
 - **ALU.A \leftarrow ID/EX.NPC**
 - If **ID/EX.SrcA == regA**
 - **ALU.A \leftarrow ID/EX.A**
 - If **ID/EX.SrcB == imm**
 - **ALU.B \leftarrow ID/EX.Imm**
 - If **ID/EX.SrcB == regB**
 - **ALU.B \leftarrow ID/EX.B**
 - Operation code:
 - **ALU.Op \leftarrow ID/EX.ALUOp**
- Storing the result:
 - **EX/MEM.ALUOut \leftarrow ALU.Out**
- Comparison unit:
 - **EX/MEM.Cond \leftarrow ID/EX.A == 0**
- Forwarding B in case of „Store” instructions
 - **EX/MEM.B \leftarrow ID/EX.B** (This will be the data to store)
- Forwarding the instruction word:
 - **EX/MEM.Instr \leftarrow ID/EX.Instr**



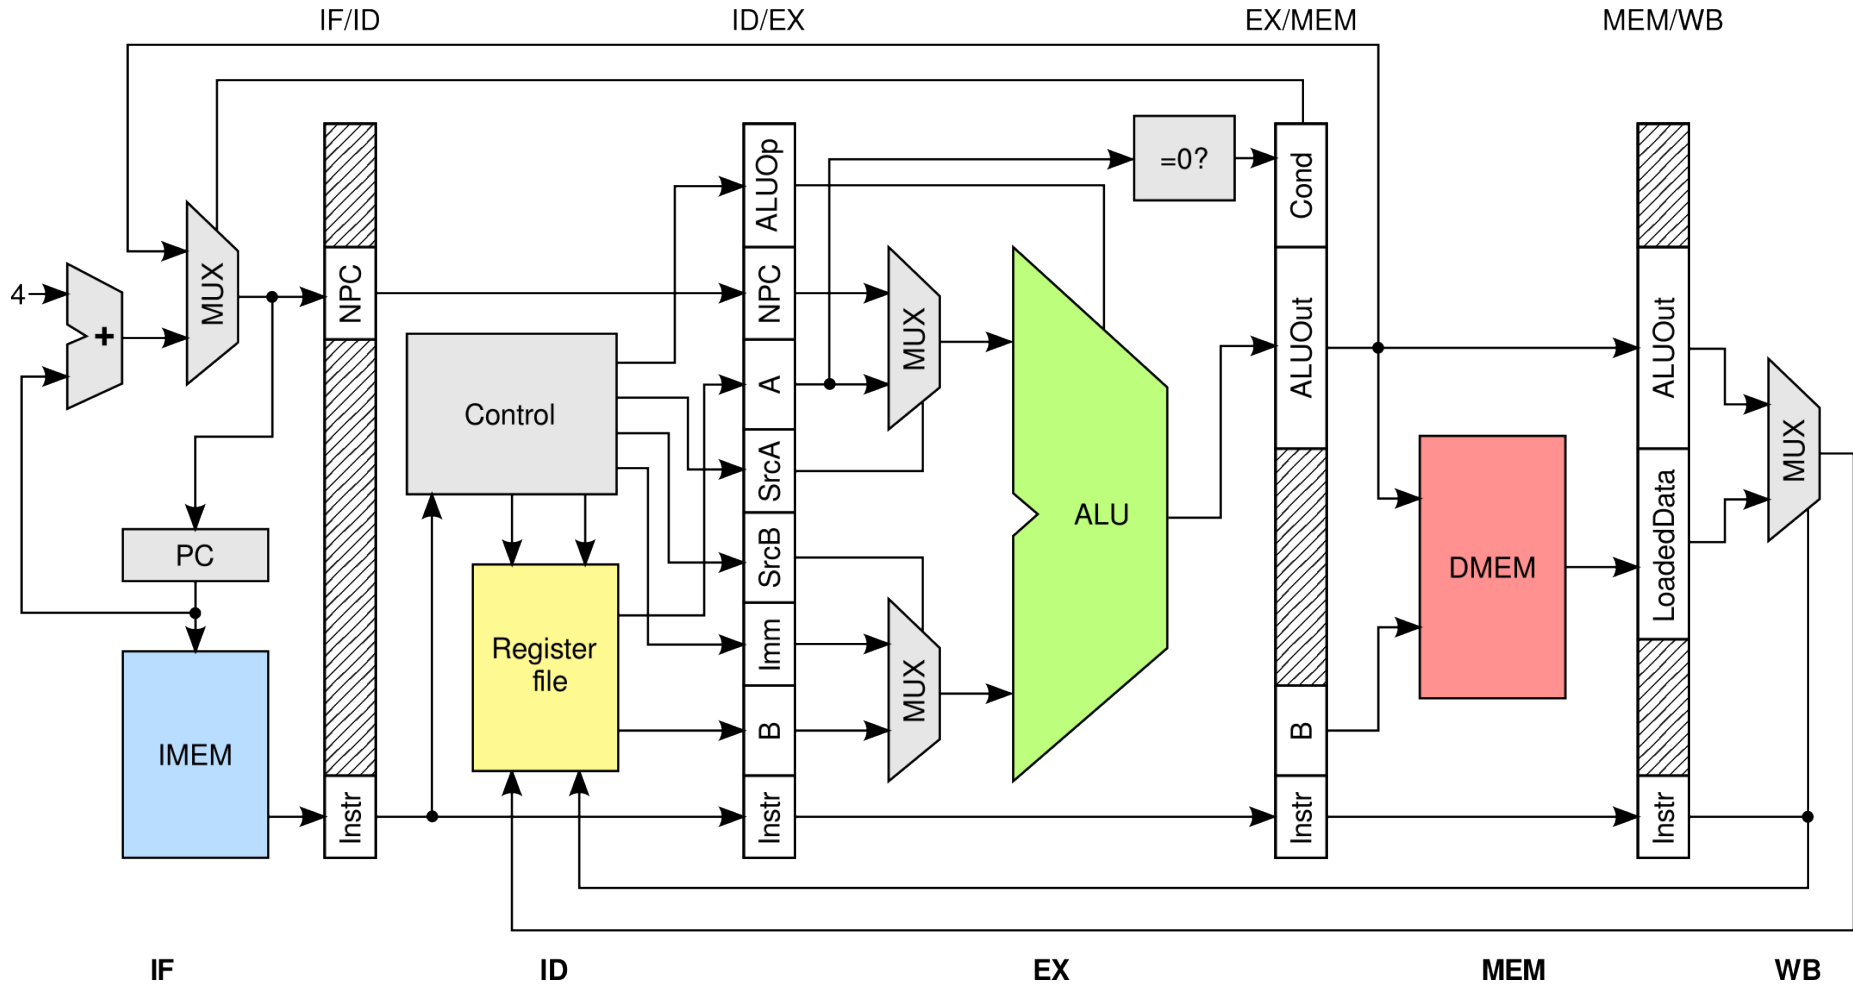
- Memory address
 - **ALUOut**
- In case of “Store” (**ID/EX.Instr.Opcod**e == Store)
 - **MEM[EX/MEM.ALUOut] ← EX/MEM.B**
(Data to write: B)
- In case of “Load” (**ID/EX.Instr.Opcod**e == Load)
 - **MEM/WB.LoadedData ← MEM[EX/MEM.ALUOut]**
- In case of arithmetic instruction:
 - **MEM/WB.ALUOut ← EX/MEM.ALUOut**
(Forwarding the result)
- Forwarding the instruction word:
 - **MEM/WB.Instr ← EX/MEM.Instr**



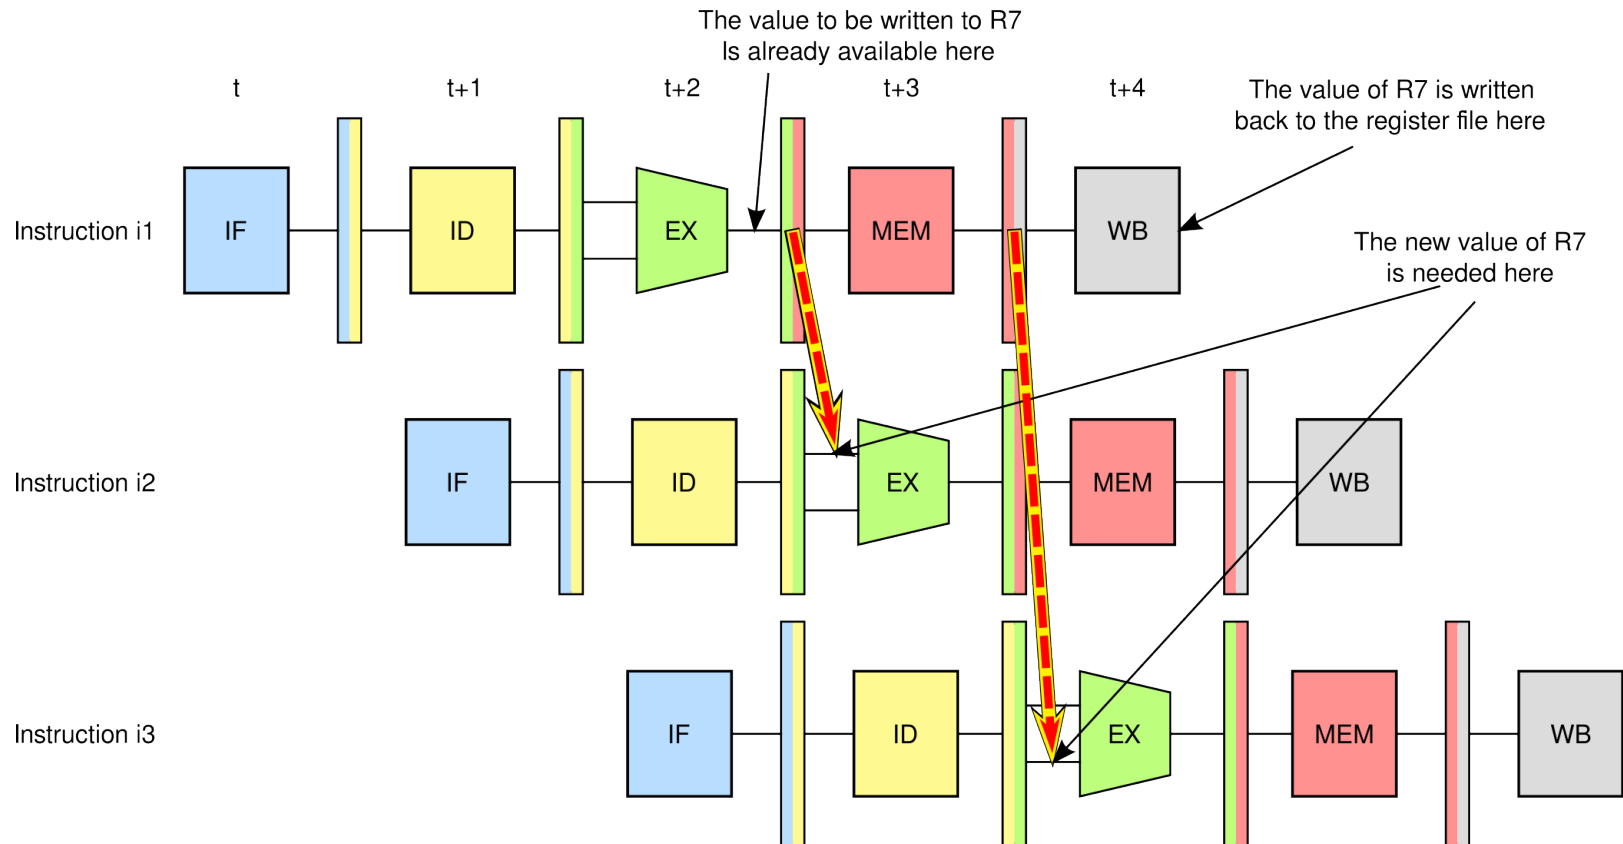
- Updating the register file (**Reg[]**)
 - In case of arithmetic instructions:
(**MEM/WB.Instr.Opcode == arithm**)
 - Reg[MEM/WB.Instr.rd] ← MEM/WB.ALUOut**
 - In case of “Load” instruction:
(**MEM/WB.Instr.Opcode == Load**)
 - Reg[MEM/WB.Instr.rd] ← MEM/WB.LoadedData**



IMPLEMENTATION OF THE PIPELINE



- Recall: we need forwarding

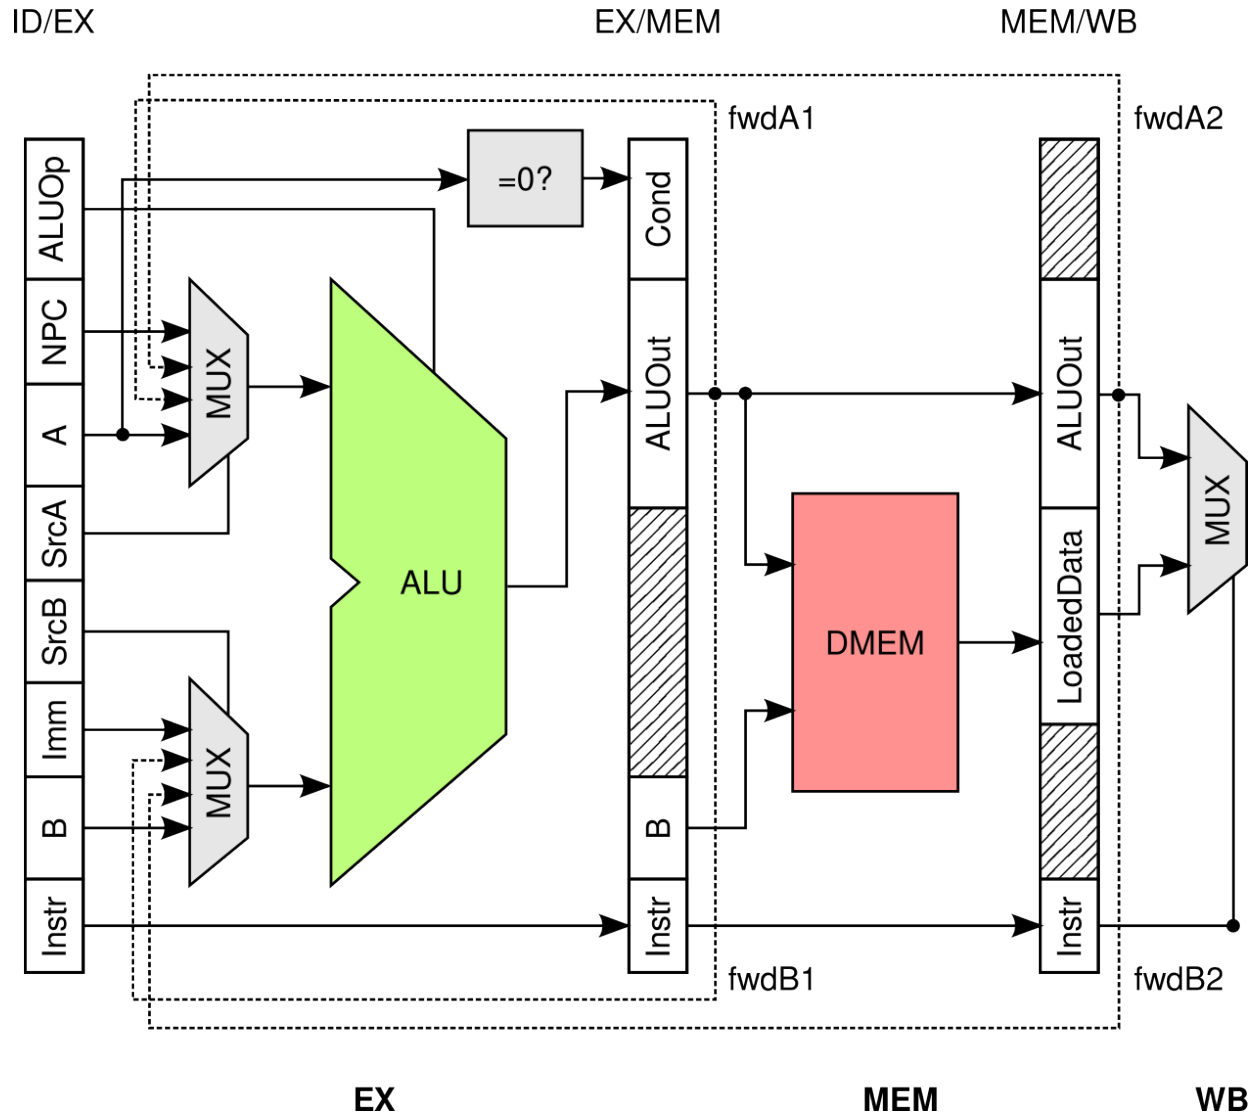


i1: R7 ← R1 + R5

i2: R8 ← R7 + R2

i3: R5 ← R8 + R7

RAW DEPENDENCIES



- How does stage ID know what the operands of ALU are?
 - There are two operands
 - Each of them can have two different sources (NPC \leftrightarrow register, constant \leftrightarrow register)
 - Register operands can come from three locations:
 - From the decoder, ID
 - From the ALU output
 - From the ALU output of the previous step
- Operand selection logic (in stage ID):
 - For example, operand 1, **ID/EX.SrcA** can be **0, 1, 2, 3**:
 - **npc**: **IF/ID.Instr.Opcod**e == **branch**
 - **fwdA1**: if **IF/ID.Instr.Opcod**e==**arithm** &&
IF/ID.Instr.ra == **ID/EX.Instr.rd**
 - **fwdA2**: if **IF/ID.Instr.Opcod**e==**arithm** &&
IF/ID.Instr.ra == **EX/MEM.Instr.rd**
 - **regA**: if **IF/ID.Instr.Opcod**e==**arithm**, otherwise
 - Similarly for operand 2

- If the resolution of a RAW dependency needs to stall the pipeline:
 - Stage ID can detect it (as it detects RAW as well)
 - Stops stage IF
- Control hazards
 - Stop & wait method: detect if a conditional jump arrived and stop stage IF for 2 cycles
 - Static prediction that always predicts „jump not taken”
 - Goes on with fetching the next instructions
 - If the jump turns out to be taken, the instructions fetched by mistake need to be invalidated:
 - Instructions are carrying a „valid” flag
 - If an instruction needs to be invalidated, „valid” is set to 0
 - Instructions with Valid=0 do nothing in stages MEM and WB (thus they leave no traces at all)

- What did we learn up to this point?
 - The concept of instruction pipelines
 - The role of pipeline registers
 - How to cope with hazards
 - How to implement an instruction pipeline
 - How to implement hazard detection and resolution
- Everything is nice till something extraordinary happens
- What can happen?
 - Exceptions!
 - An interrupt from an I/O device
 - Page fault
 - Protection fault
 - Invalid instruction
 - Etc.

- Ideal behavior
 - Given that the exception raises during instruction i .
 - We want to ensure a state where
 - All instructions before i are completed
 - Instructions after i have not been started at all
- ***precise exceptions***
- Not easy to implement if we have a pipeline
- When an exception happens, we want to prevent the started but not yet completed instructions from leaving a trace
- ... but sometimes it can not be prevented
(a later instruction might have already changed the content of the memory)

- What can happen in the various stages:
 - Stage IF: page fault, protection fault
 - Stage ID: invalid instruction
 - Stage EX: arithmetic error (integer overflow, division by zero, etc.)
 - Stage MEM: page fault, protection fault
 - Stage WB: no exceptions can happen

	1	2	3	4	5	6
$R_k \leftarrow R_m + R_n$	IF	ID	EX	MEM	WB	
$R_i \leftarrow \text{MEM}[R_j]$		IF	ID	EX	MEM	WB

- Instruction 1.: Integer overflow in stage EX
- Instruction 2.: page fault in stage IF
→ The order of events is not correct! The exception of a later instruction occurs sooner!

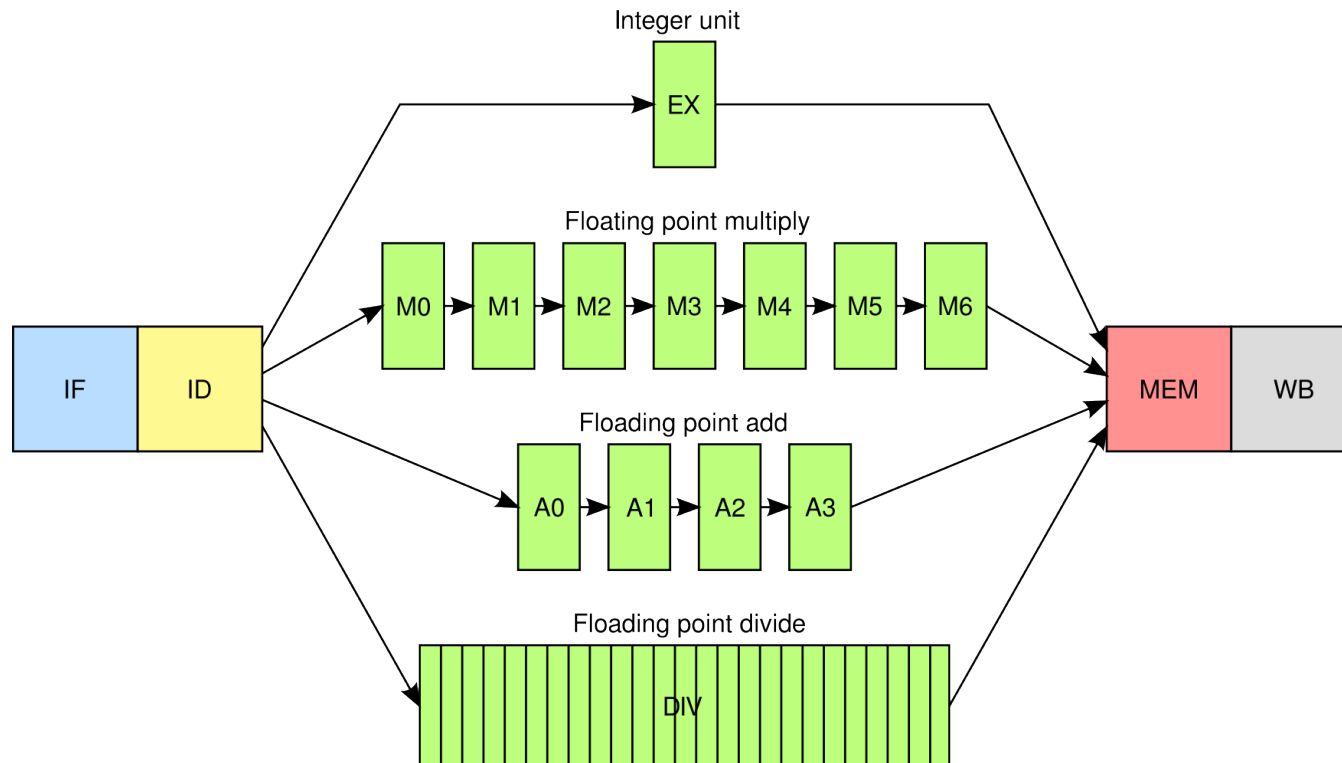
- A possible solution:
 - Exceptions are not handled immediately
 - The CPU just marks the instruction with a flag
 - Handling of exception: only in stage WB
 - the order of exceptions will be correct this way

	1	2	3	4	5	6
$R_k \leftarrow R_m + R_n$	IF	ID	EX	MEM	WB	
$R_i \leftarrow \text{MEM}[R_j]$		IF	ID	EX	MEM	WB

- Until now all arithmetic operations were executed in 1 cycle
 - But different arithmetic operations may have different **latencies**
 - Floating point operations need more cycles
 - Multiplication needs more cycles than addition
 - Division is the slowest operation
 - The arithmetic functional units can be pipelined or unpipelined internally
 - E.g. the floating point addition might have a latency of 4 cycles,
 - ...but can start a new addition operation every cycle
- **Iteration interval** = 1

PIPELINE WITH MULTICYCLE OPERATIONS

Functional unit	Latency	Iteration interval
Integer ALU	1	1
FP add	4	1
FP multiply	7	1
FP divide	25	25



- Observe:
 - Instructions are started *in-order*, but completed *out-of-order*

D4 ← D1 * D5	IF	ID	M0	M1	M2	M3	M4	M5	M6	MEM	WB
D2 ← D1 + D3		IF	ID	A0	A1	A2	A3	MEM	WB		
D0 ← MEM [R0+4]			IF	ID	EX	MEM	WB				
MEM [R0+8] ← D5				IF	ID	EX	MEM	WB			

- In this example this was not a problem
- ...but it can introduce problems in other situations
- Stage ID has new tasks:
 - To keep the semantics of the program
 - To cope with new kinds of hazards and dependencies

STRUCTURAL HAZARDS INTRODUCED

Instruction	1	2	3	4	5	6	7	8	9	10	11
$D4 \leftarrow D1 * D5$	IF	ID	M0	M1	M2	M3	M4	M5	M6	MEM	WB
...		IF	ID	EX	MEM	WB					
...			IF	ID	EX	MEM	WB				
$D2 \leftarrow D1 + D3$				IF	ID	A0	A1	A2	A3	MEM	WB
...					IF	ID	EX	MEM	WB		
...						IF	ID	EX	MEM	WB	
$D0 \leftarrow \text{MEM}[R0+4]$							IF	ID	EX	MEM	WB

- Problem:
 - 10.: there are 3 instructions in stage MEM at the same time
 - 11.: there are 3 instructions in stage WB at the same time
- Solution:
 - Stage ID knows exactly the latency of the instructions → it can detect such situations
 - If ID thinks that an instruction will cause structural hazard later → it does not let it move forward in the pipeline (stall)

- Data hazards are worse than before
- Problem: pipeline is longer
 - The result of an operation is calculated late
→ if an instruction uses that result, it has to wait a lot
 - There are more instructions under execution at the same time
→ the chance of having a RAW dependency among them is larger (compared to our simple 5-stage pipeline)

Instruction	1	2	3	4	5	6	7	8	9	10	11	12	14	14	15	16	17
D4 ← MEM [R0+4]	IF	ID	EX	ME	WB												
D0 ← D4 * D6		IF	ID	Stop	M0	M1	M2	M3	M4	M5	M6	ME	WB				
D2 ← D0 + D8			IF	Stop	ID	Stop	Stop	Stop	Stop	Stop	Stop	A0	A1	A2	A3	ME	WB
MEM [R0+4] ← D2					IF	Stop	Stop	Stop	Stop	Stop	Stop	ID	EX	Stop	Stop	Stop	ME

- WAW dependency:
 - Two instructions write their results to the same register
- If instructions are allowed to complete out-of-order → might introduce problems!
 - It might happen that the later instruction writes its result to the target register sooner
 - Stage ID has to detect this situation and handle it by inserting stalls

Instruction	1	2	3	4	5	6	7	8
D2 ← D1 + D3	IF	ID	A0	A1	A2	A3	MEM	WB
...		IF	ID	EX	MEM	WB		
D2 ← MEM [R0+4]			IF	ID	EX	MEM	WB	

- Clearly bad
 - Solution: Stage ID introduces 2 stalls at the last instruction

- Instruction pipeline is a nice, beautiful concept
 - Difficulties:
 - Hazards / data dependencies
 - Exceptions
 - Multi-cycle operations
- The longer the pipeline is, the more serious these difficulties are